

Nathaniel Oliver
nathanielsoliver@gmail.com | [Portfolio](#)

May 28, 2026

Hiring Team OCBridge San Jose, CA

Dear OCBridge Hiring Team,

I'm applying for the Full Stack Engineering Intern role. OCBridge is building AI-native recruiting infrastructure — and it happens to be the exact problem space I've been building in. My most substantial project is an [end-to-end agentic resume tailoring system](#): a local-first Python application that ingests resume, GitHub, and LinkedIn data into a knowledge graph, then uses multi-agent LangGraph workflows and the OpenAI and Anthropic APIs to generate evidence-backed resume improvements and ATS gap analysis. I architected the full stack myself — PostgreSQL backend, a service layer decoupling database access from LLM orchestration, ETL pipelines for document ingestion, and a CLI/TUI interface — and shipped it with pytest coverage across service boundaries and agent routing. Building that system gave me direct experience with the same problems your team is solving: automating talent workflows with LLMs, making AI outputs reliable enough for production use, and structuring backend systems that can actually scale.

My production experience reinforces the backend and data pipeline side. As a Data Science Intern at IDX Exchange, I built a [live ML pipeline](#) processing 100K+ real estate records — handling geospatial filtering, custom imputation, and automated hyperparameter optimization — and deployed a stacked ensemble model (XGBoost + LightGBM) achieving accuracy on par with Redfin's valuation model. That internship taught me what it takes to move from prototype to production: data quality checks, reproducible pipelines, and systems that behave consistently under real-world variation.

On the frontend, I built an [interactive web application](#) in JavaScript and D3.js for UCSD Biomedical Research, delivering dynamic dashboards for exploring high-resolution circadian time-series data from mouse experiments. My [portfolio site](#) goes further: it's a vanilla JS ES module app with a custom component architecture, autocomplete search with keyboard navigation, and a Cloudflare Worker serving as a REST backend — automated end-to-end with Node.js scripts that sync project data from the GitHub API.

I'm drawn to OCBridge specifically because your platform sits at the intersection of AI agent workflows and a domain where reliability and trust matter — a recruiter acting on Hiring Copilot's output needs to trust it. I've spent time thinking carefully about that exact tradeoff in my own work, and I'd like to bring that perspective to your team. I'm fast to pick up new tools, comfortable with ambiguity, and most productive when the work is real and the feedback loop is short — which is exactly what you're describing.

I'd love to talk.

Nathaniel Oliver
nathanielsoliver@gmail.com